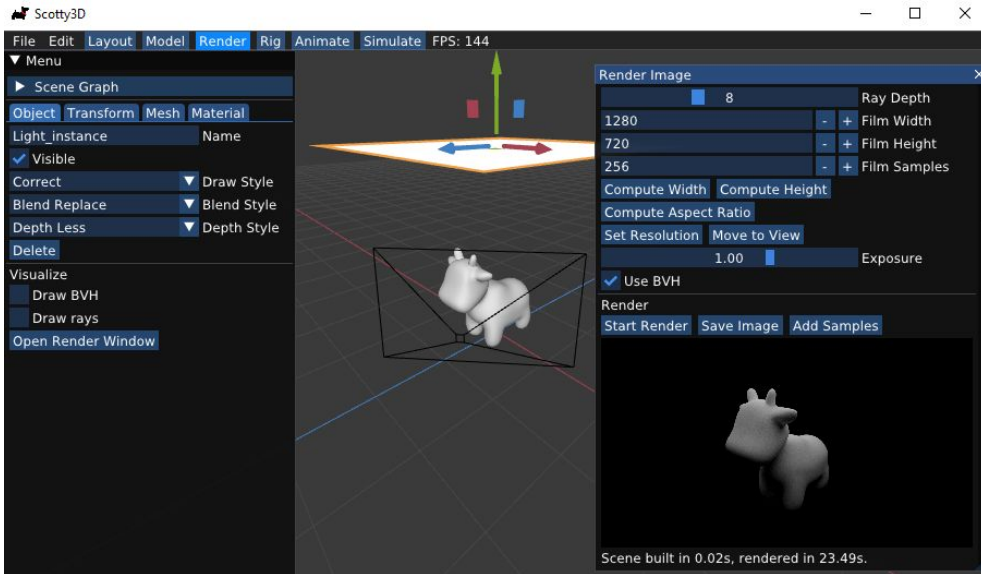


# Scotty3D

- Custom graphics package used in 15462/662 Computer Graphics; C++
- Supports rasterization, raytracing, mesh editing, and particle simulation



# Halfedge Mesh Data Structure

- Local connectivity description with explicit lists of faces, edges, vertices, and 'halfedges'

```
struct Halfedge
```

```
{
```

```
    Halfedge* twin;
```


```
    Halfedge* next;
```

```
    Vertex* vertex;
```

```
    Edge* edge;
```


```
    Face* face;
```

```
};
```



# Polygon Triangulation



- Input : Any manifold\* mesh
- Output : Same mesh with every face degree = 3
- Loop through every face input mesh has and subdivide it into triangles
- New face / edges / halfedges created in the process => Problem for parallelization!



Manifold = “can exist in real world”

# Serial Solution

- Loop through every face, perform fan triangulation
  - Fan triangulation : Choose a vertex and draw edges to all other vertices in the face
- Works regardless of concave / convex
  - Assume inverted faces are ok
- Let a face have  $n$  vertices, this generates  $n-3$  new faces,  $n-3$  new edges, and  $2*(n-3)$  new halfedges
- **Problem** : we have temporal locality but no spatial locality  
=> a lot of memory lookups & cache misses
  - Parallelize to hide latency of memory lookups!



# Difficulties in Parallelization

## 1. Very hard to predict workload

- a. Unknown distribution of face degrees until we inspect each face

## 2. Cannot benefit from memory coherency

- a. No elements consisting a face is guaranteed to be contiguous in memory
- b. No two faces of adjacent indices are guaranteed to have contiguous/nearby elements

## 3. Halfedge Mesh in Scotty3D uses lists to store elements

- a. If we use shared memory (OpenMP), no built-in thread-safe way to emplace new elements
- b. If we use isolated memory (Open MPI), iterators used to store references to new elements will be invalid once merged with other processes

# Parallel Solution 1 (OpenMP)

- Implement thread-safe custom  
emplace functions
  - Use C++'s `std::mutex`  
(synchronization primitive for  
exclusive, non-recursive ownership)
  - Each mesh has a unique pointer to  
a mutex instance

=> Lock mutex when a thread-safe  
emplace gets called, unlock when  
done
- `#pragma omp parallel for`
  - Experiment with various scheduling  
options (static, dynamic, guided)
  - Experiment with parallelization  
conditions (if `faces.size() <`  
threshold)
- Dynamically allocated objects  
(mesh elements) are in shared  
memory by default

# Parallel Solution 2 (Open MPI)

- Change lists to vectors


- Instead of storing references to iterators within lists, store indices into vectors

- Three-stage implementation

- Distribution : MPI\_Bcst to send relevant mesh structure to each process
- Computation : Perform fan triangulation, creating new faces, edges, and halfedges
- Collection : New mesh elements are appended to the end of their respective element vectors


Broadcasting entire mesh

Normalized Time Profile for Each Stage



Broadcasting relevant parts only

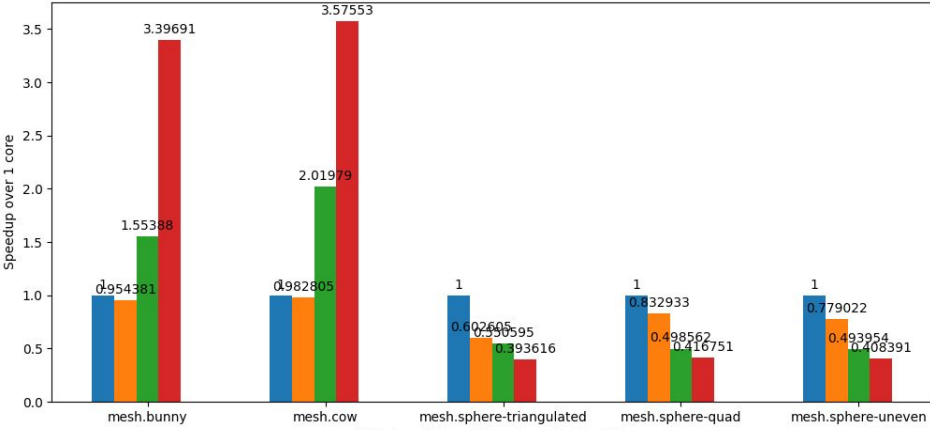
Normalized Time Profile for Each Stage




- Ensure connectivity modifications are propagated when merged
- Ensure any index changes due to previously merged elements is reflected to new connections

# Results - OpenMP


OMP for Speedup with Mac M1



OMP guided schedule Speedup with Mac M1



OMP shared (faces) Speedup with Mac M1




--- OMP Time Table ---

benchmark	mesh.bunny	mesh.cow	mesh.sphere-triangulated	mesh.sphere-quad	mesh.sphere-uneven
1	0.444899	0.642631	9.8455e-05	0.000949383	0.00130997
2	0.382854	0.53734	0.000148571	0.00112353	0.000960921
4	0.253271	0.355731	0.000168143	0.00143004	0.00101812
8	0.135276	0.188273	0.000236651	0.00148555	0.0011528
16	0.0793805	0.119626	0.000437191	0.00179665	0.00149899

--- OMP Speedup Table ---

benchmark	mesh.bunny	mesh.cow	mesh.sphere-triangulated	mesh.sphere-quad	mesh.sphere-uneven
1	1.0	1.0	1.0	1.0	1.0
2	1.16205917	1.19594856	0.6626798	0.84500013	1.36324422
4	1.75661248	1.80650829	0.61479428	0.66388563	1.2866558
8	3.28882433	3.41329346	0.41603458	0.63907846	1.13633761
16	5.6102925	5.37200107	0.22519985	0.52841845	0.87859074


# Results - OpenMPI



Performance counter stats for 'mpirun -np 2 ./Scotty3D --run-tests a2.g1.mpi.triangulate.mesh.bunny' (20 runs):

877.43 msec task-clock	#	1.029 CPUs utilized	(+ 0.64%)
3285 context-switches	#	3.777 K/sec	(+ 0.46%)
24 cpu-migrations	#	27.593 /sec	(+ 1.26%)
33317 page-faults	#	38.305 K/sec	(+ 0.01%)
3501146801 cycles	#	4.025 GHz	(+ 0.36%) (62.42%)
3854856292 instructions	#	1.08 insn per cycle	(+ 0.61%) (72.48%)
862829741 branches	#	991.996 M/sec	(+ 0.37%) (72.53%)
18077366 branch-misses	#	1.17% of all branches	(+ 1.67%) (72.13%)
1179625844 L1-dcache-loads	#	1.356 G/sec	(+ 0.44%) (71.93%)
74698036 L1-dcache-load-misses	#	6.47% of all L1-dcache accesses	(+ 0.47%) (71.13%)
18463828 LLC-loads	#	21.228 M/sec	(+ 0.36%) (70.19%)
5590263 LLC-load-misses	#	30.35% of all LL-cache accesses	(+ 0.90%) (70.70%)
<not supported> L1-icache-loads			
10990290 L1-icache-load-misses			(+ 1.13%) (70.29%)
1068389970 dTLB-loads	#	1.228 G/sec	(+ 0.59%) (69.83%)
1796606 dTLB-load-misses	#	0.17% of all dTLB cache accesses	(+ 1.06%) (59.29%)
84488 iTLB-loads	#	97.136 K/sec	(+ 3.29%) (60.55%)
113974 iTLB-load-misses	#	145.49% of all iTLB cache accesses	(+ 6.07%) (60.55%)
<not supported> L1-dcache-prefetches			
<not supported> L1-dcache-prefetch-misses			

0.85283 +- 0.00376 seconds time elapsed (+ 0.44%)



Performance counter stats for 'mpirun -np 8 ./Scotty3D --run-tests a2.g1.mpi.triangulate.mesh.bunny' (20 runs):

4011.15 msec task-clock	#	3.681 CPUs utilized	(+ 0.29%)
16681 context-switches	#	4.083 K/sec	(+ 0.25%)
582 cpu-migrations	#	142.444 /sec	(+ 1.74%)
97236 page-faults	#	23.798 K/sec	(+ 0.01%)
14651647235 cycles	#	3.586 GHz	(+ 0.17%) (62.96%)
12528901141 instructions	#	0.86 insn per cycle	(+ 0.37%) (70.71%)
2632642215 branches	#	644.335 M/sec	(+ 0.23%) (71.03%)
21520746 branch-misses	#	0.81% of all branches	(+ 0.62%) (71.10%)
3707856892 L1-dcache-loads	#	907.493 M/sec	(+ 0.30%) (71.24%)
186499801 L1-dcache-load-misses	#	4.91% of all L1-dcache accesses	(+ 1.12%) (70.88%)
55505640 LLC-loads	#	13.585 M/sec	(+ 0.82%) (70.01%)
29875168 LLC-load-misses	#	54.02% of all LL-cache accesses	(+ 0.23%) (69.27%)
<not supported> L1-icache-loads			
20261142 L1-icache-load-misses			(+ 3.40%) (67.98%)
3653327974 dTLB-loads	#	894.147 M/sec	(+ 0.53%) (67.66%)
4289576 dTLB-load-misses	#	0.12% of all dTLB cache accesses	(+ 0.49%) (60.41%)
255151 iTLB-loads	#	62.448 K/sec	(+ 3.78%) (61.33%)
677184 iTLB-load-misses	#	230.46% of all iTLB cache accesses	(+ 3.15%) (62.50%)
<not supported> L1-dcache-prefetches			
<not supported> L1-dcache-prefetch-misses			

1.08956 +- 0.00403 seconds time elapsed (+ 0.37%)